

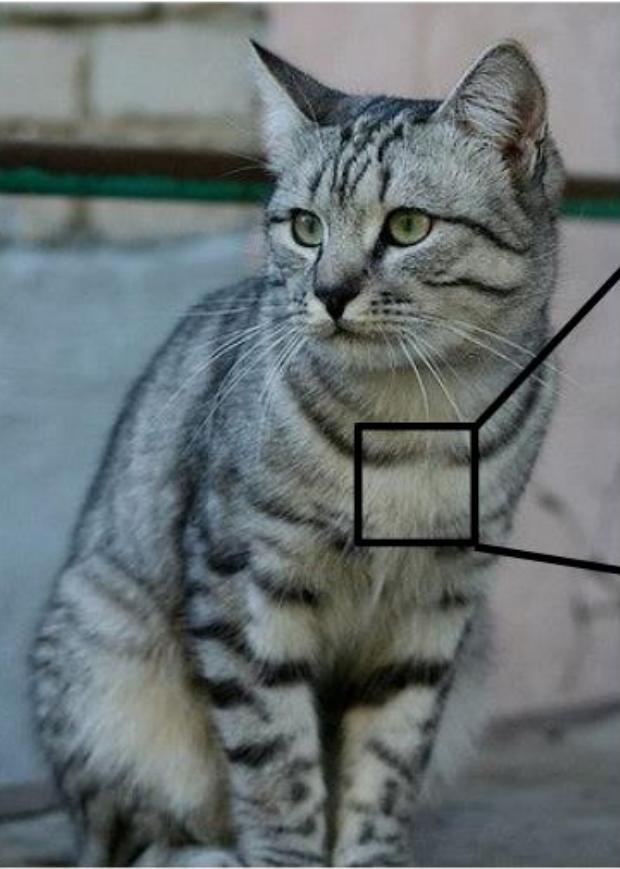


Introduction to Computer Vision

24/09/2019

Computer
Vision is
Challenging





What we see

```
[[105 112 108 111 104 99 106 99 96 103 112 119 104 97 93 87]
 [91 98 102 106 104 79 98 103 99 105 123 136 110 105 94 85]
 [76 85 90 105 128 105 87 96 95 99 115 112 106 103 99 85]
 [99 81 81 93 120 131 127 100 95 98 102 99 96 93 101 94]
 [106 91 61 64 69 91 88 85 101 107 109 98 75 84 96 95]
 [114 108 85 55 55 69 64 54 64 87 112 129 98 74 84 91]
 [133 137 147 103 65 81 80 65 52 54 74 84 102 93 85 82]
 [128 137 144 140 109 95 86 70 62 65 63 63 60 73 86 101]
 [125 133 148 137 119 121 117 94 65 79 80 65 54 64 72 98]
 [127 125 131 147 133 127 126 131 111 96 89 75 61 64 72 84]
 [115 114 109 123 150 148 131 118 113 109 100 92 74 65 72 78]
 [89 93 90 97 188 147 131 118 113 114 113 109 106 95 77 80]
 [63 77 86 81 77 79 102 123 117 115 117 125 125 130 115 87]
 [62 65 82 89 78 71 80 101 124 126 119 101 107 114 131 119]
 [63 65 75 88 89 71 62 81 120 138 135 105 81 98 110 118]
 [87 65 71 87 106 95 69 45 76 130 126 107 92 94 105 112]
 [118 97 82 86 117 123 116 66 41 51 95 93 89 95 102 107]
 [164 146 112 80 82 120 124 104 76 48 45 66 88 101 102 109]
 [157 170 157 120 93 86 114 132 112 97 69 55 70 82 99 94]
 [130 128 134 161 139 100 109 118 121 134 114 87 65 53 69 86]
 [128 112 96 117 150 144 120 115 104 107 102 93 87 81 72 79]
 [123 107 96 86 83 112 153 149 122 109 104 75 80 107 112 99]
 [122 121 102 80 82 86 94 117 145 148 153 102 58 78 92 107]
 [122 164 148 103 71 56 78 83 93 103 119 139 102 61 69 84]]]
```

What a computer sees

Challenges: Background Clutter



[This image is CC0 1.0 public domain](#)



[This image is CC0 1.0 public domain](#)

Challenges: Illumination



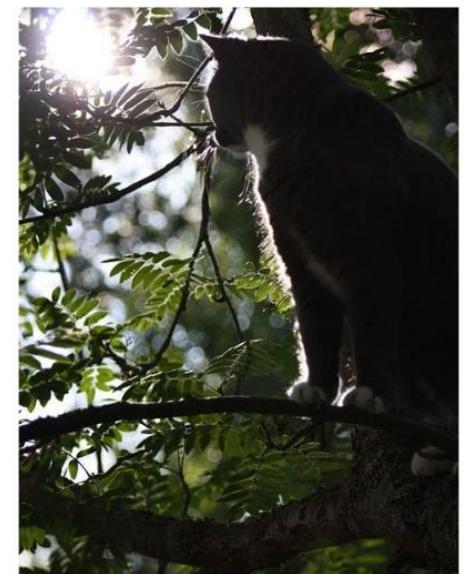
[This image is CC0 1.0 public domain](#)



[This image is CC0 1.0 public domain](#)

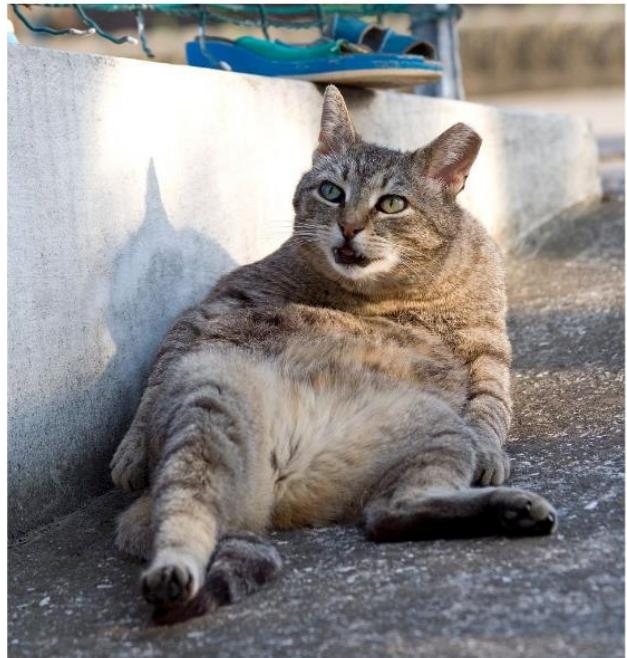


[This image is CC0 1.0 public domain](#)



[This image is CC0 1.0 public domain](#)

Challenges: Deformation



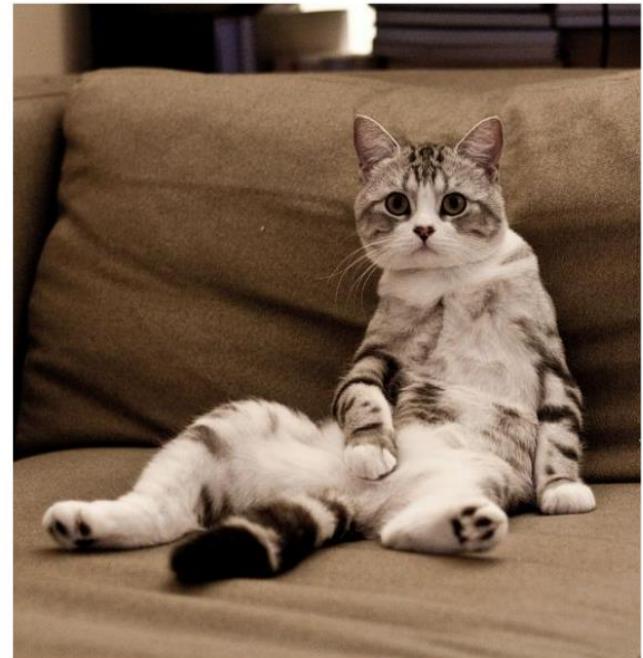
[This image by Umberto Salvagnin](#)
is licensed under [CC-BY 2.0](#)



[This image by Umberto Salvagnin](#)
is licensed under [CC-BY 2.0](#)



[This image by sare bear is](#)
licensed under [CC-BY 2.0](#)



[This image by Tom Thai is](#)
licensed under [CC-BY 2.0](#)

Challenges: Occlusion



[This image is CC0 1.0 public domain](#)



[This image is CC0 1.0 public domain](#)



[This image by **jonsson** is licensed under CC-BY 2.0](#)

Challenges: Intraclass variation



This image is [CC0 1.0 public domain](#)

dog



mug



hat





How do we solve it?

Deep Learning

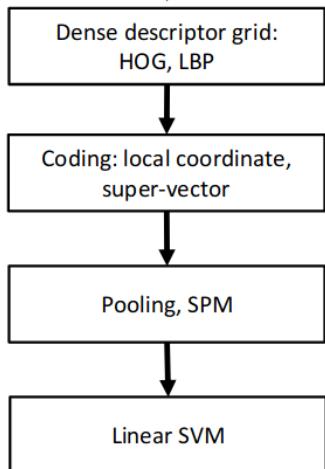


Deep Convolutional Neural Networks
have become an important tool for Computer Vision.

IMAGENET Large Scale Visual Recognition Challenge

Year 2010

NEC-UIUC

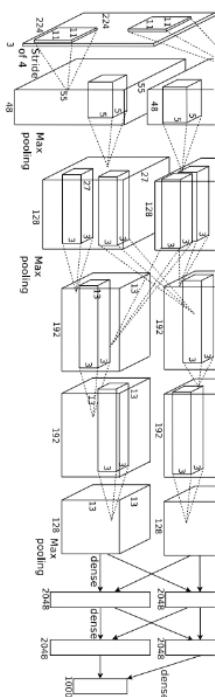


[Lin CVPR 2011]

Lion image by Swissfrog is licensed under CC BY 3.0

Year 2012

SuperVision



[Krizhevsky NIPS 2012]

Figure copyright Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton, 2012. Reproduced with permission.

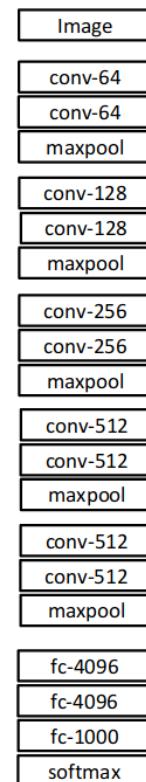
Year 2014

GoogLeNet

- Pooling
- Convolution
- Softmax
- Other



VGG

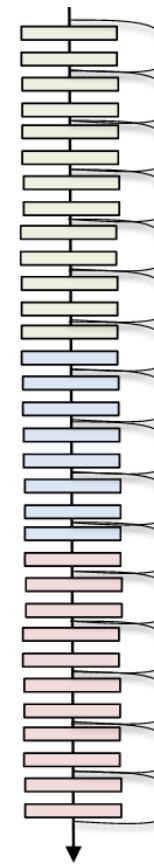


[Szegedy arxiv 2014]

[Simonyan arxiv 2014]

Year 2015

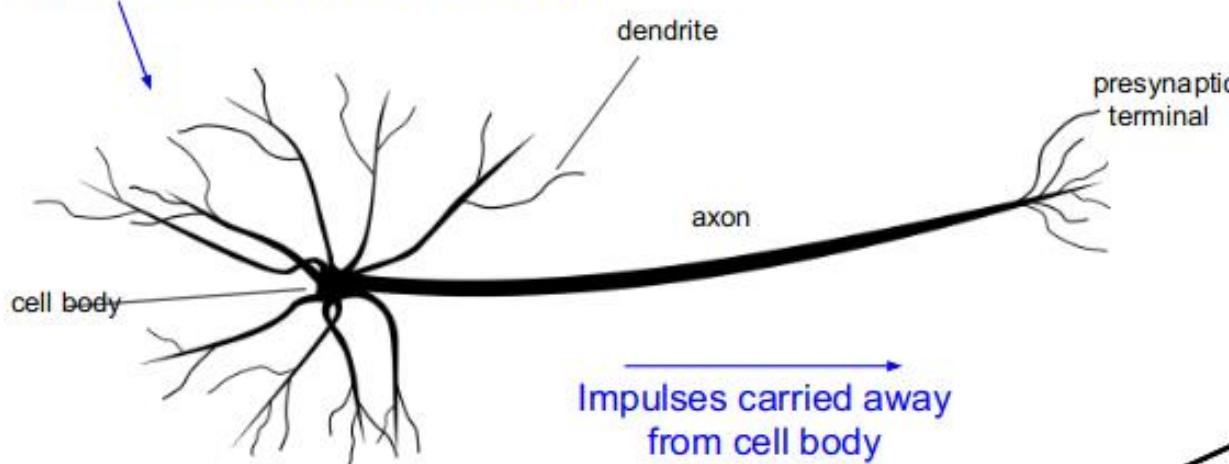
MSRA



[He ICCV 2015]

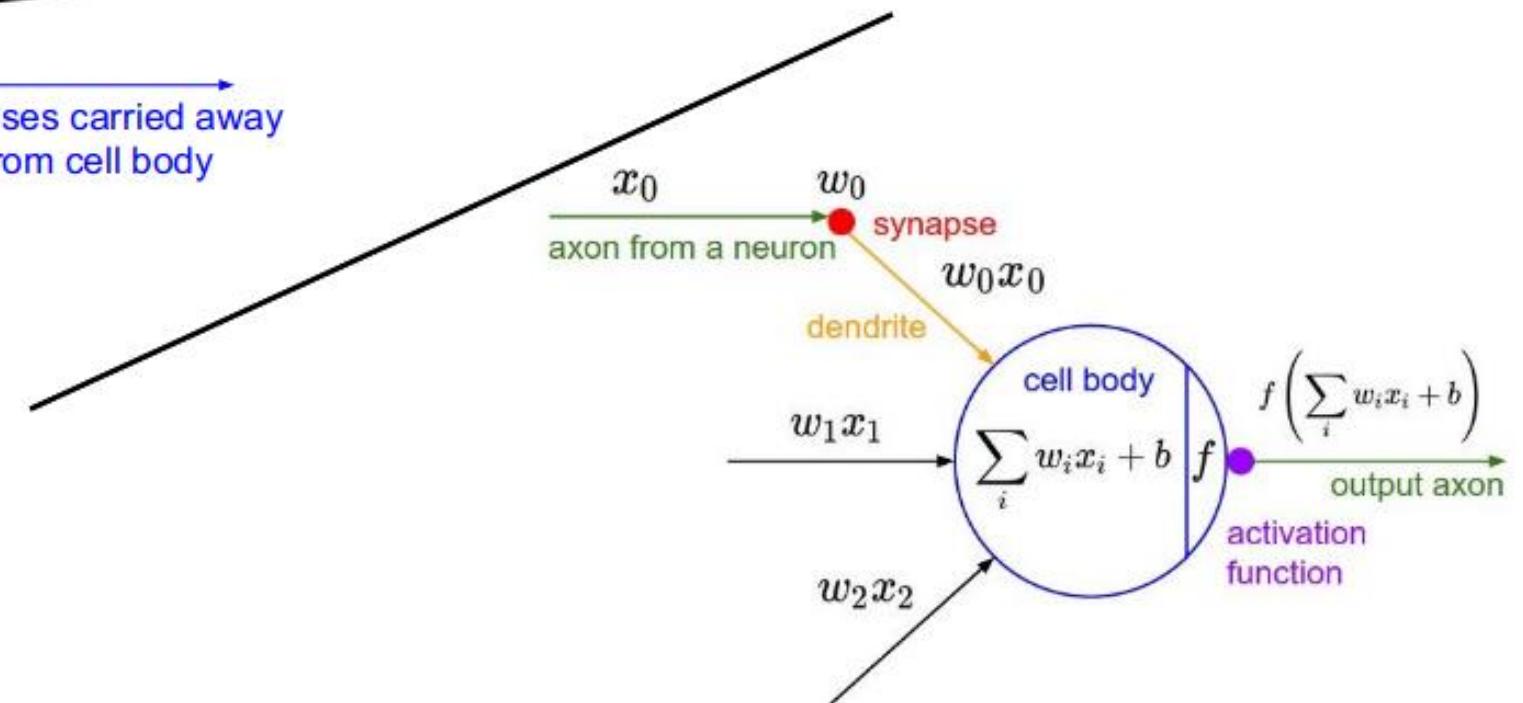
Neural Networks

Impulses carried toward cell body



This image by Felipe Perucho
is licensed under CC-BY 3.0

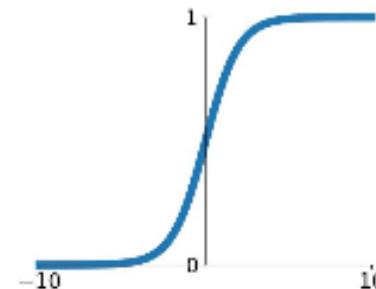
Impulses carried away
from cell body



Activation Functions: Non-linear Computations

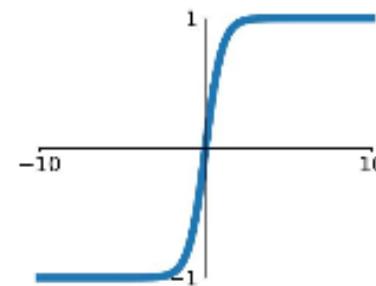
Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



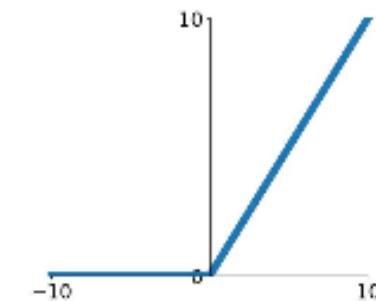
tanh

$$\tanh(x)$$



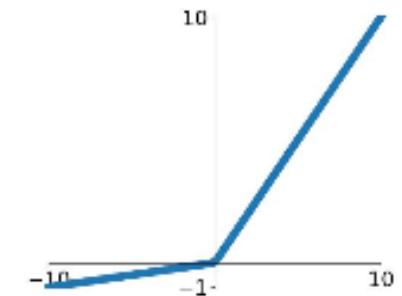
ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$

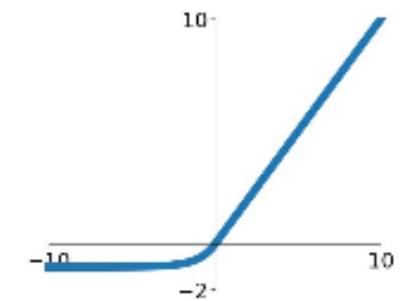


Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



Neural Networks

- Linear score function

$$f = Wx$$

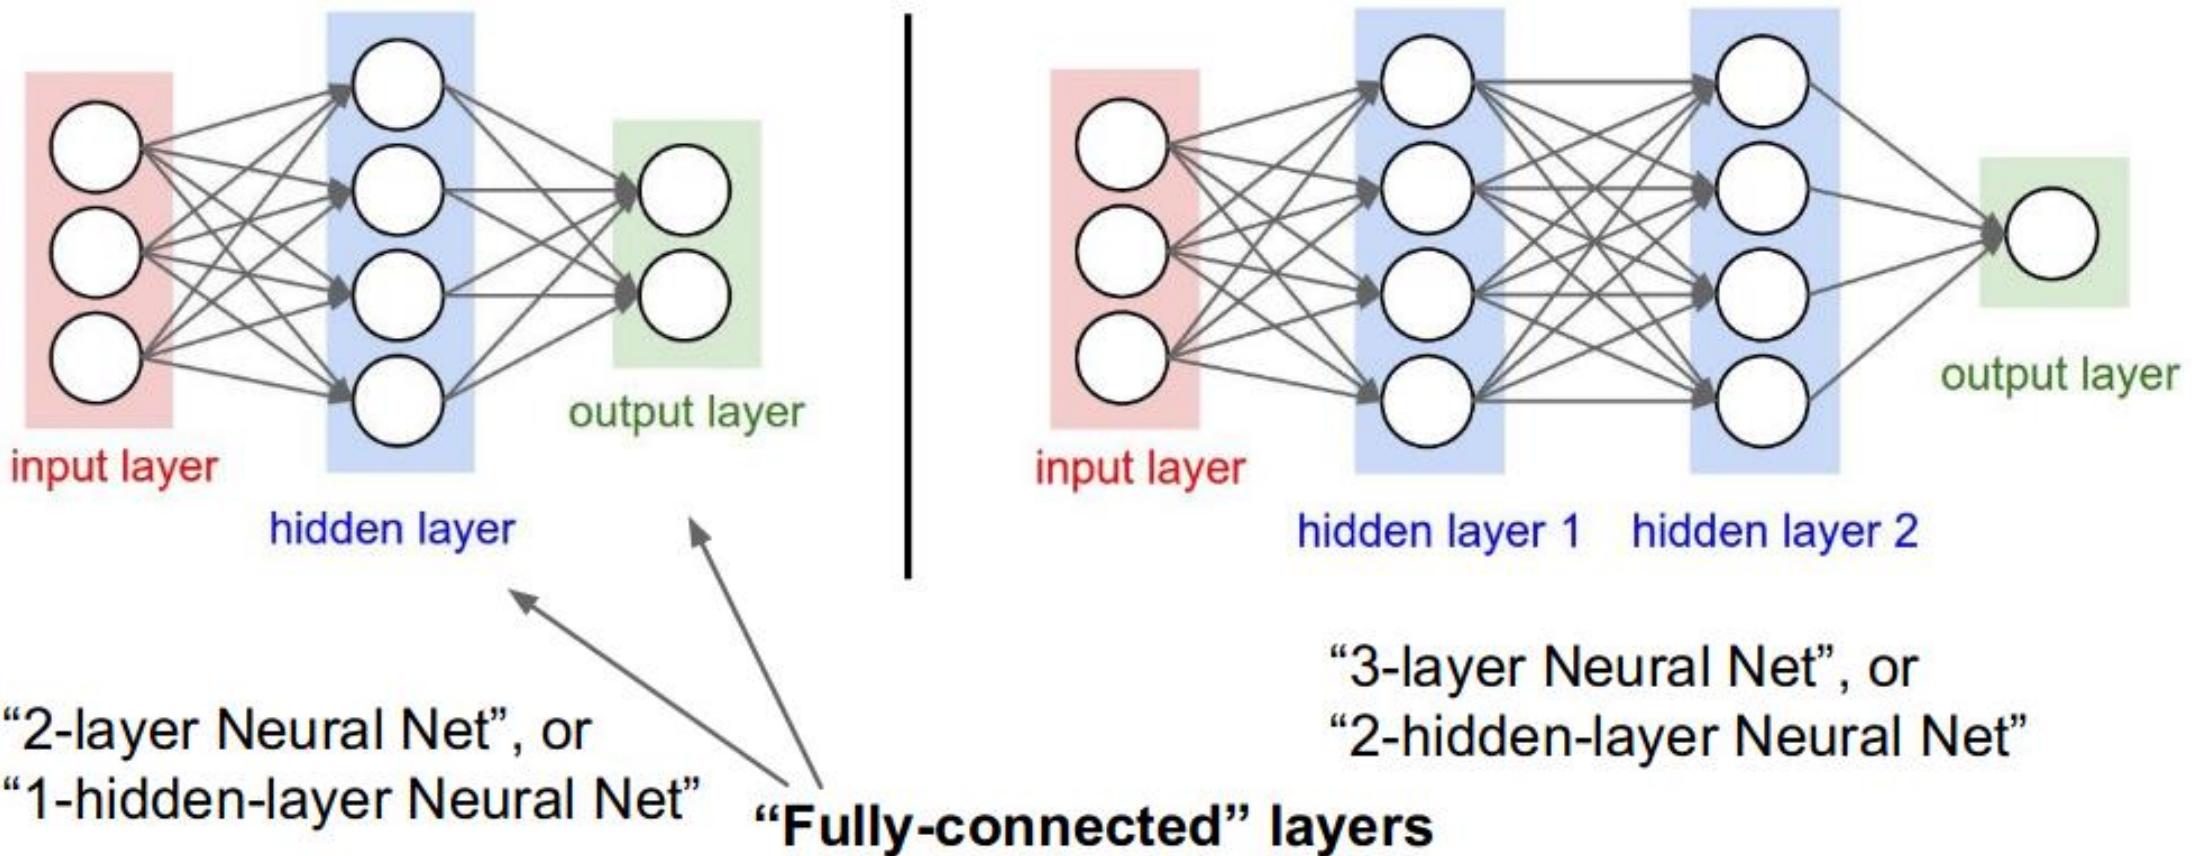
- 2-layer Neural Network

$$f = W_2 \max(0, W_1 x)$$

- 3-layer Neural Network

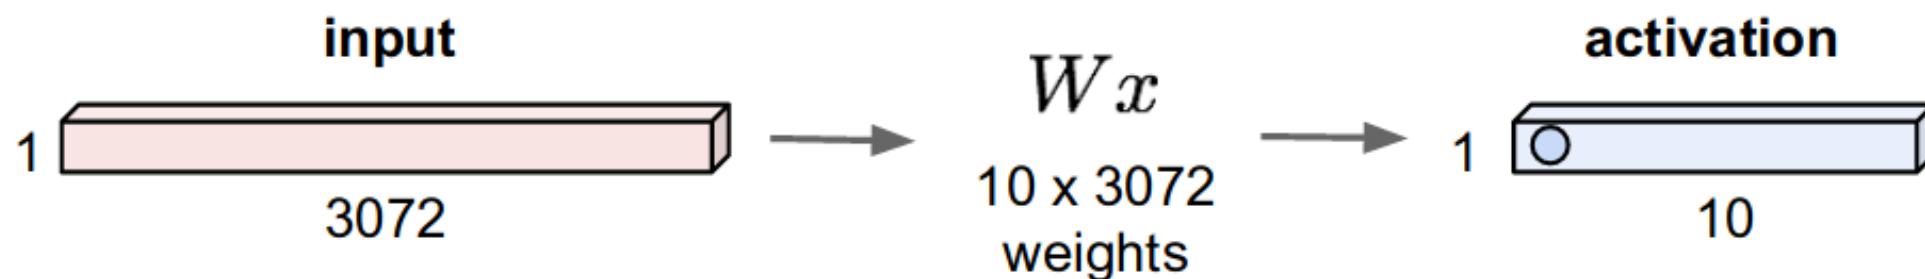
$$f = W_3 \max(0, W_2 \max(0, W_1 x))$$

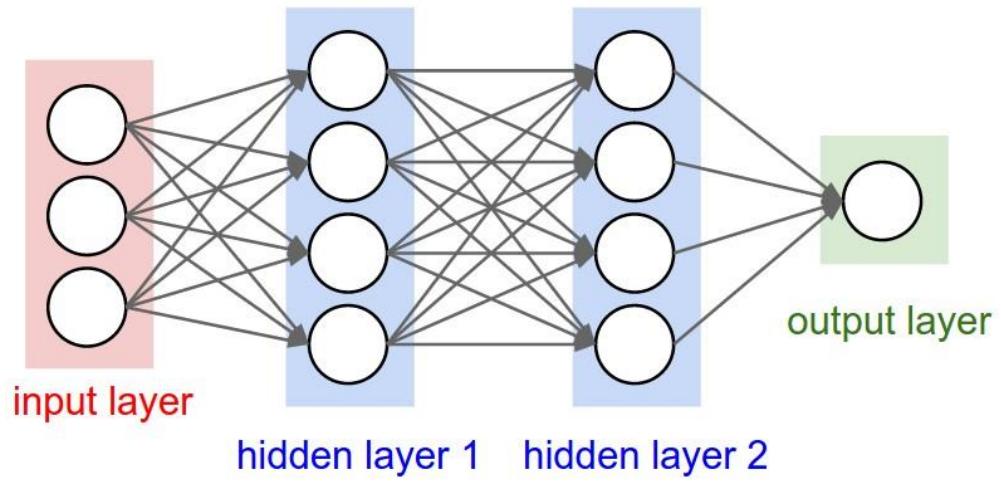
Neural Networks



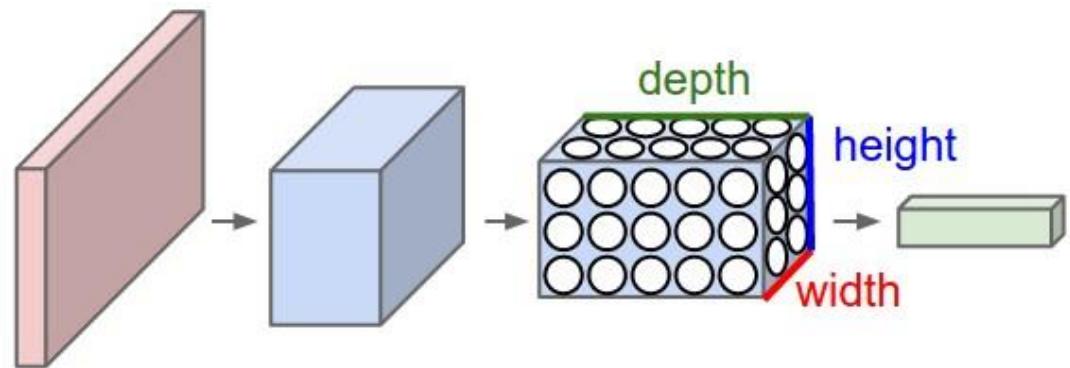
Fully Connected Layer

32x32x3 image -> stretch to 3072 x 1



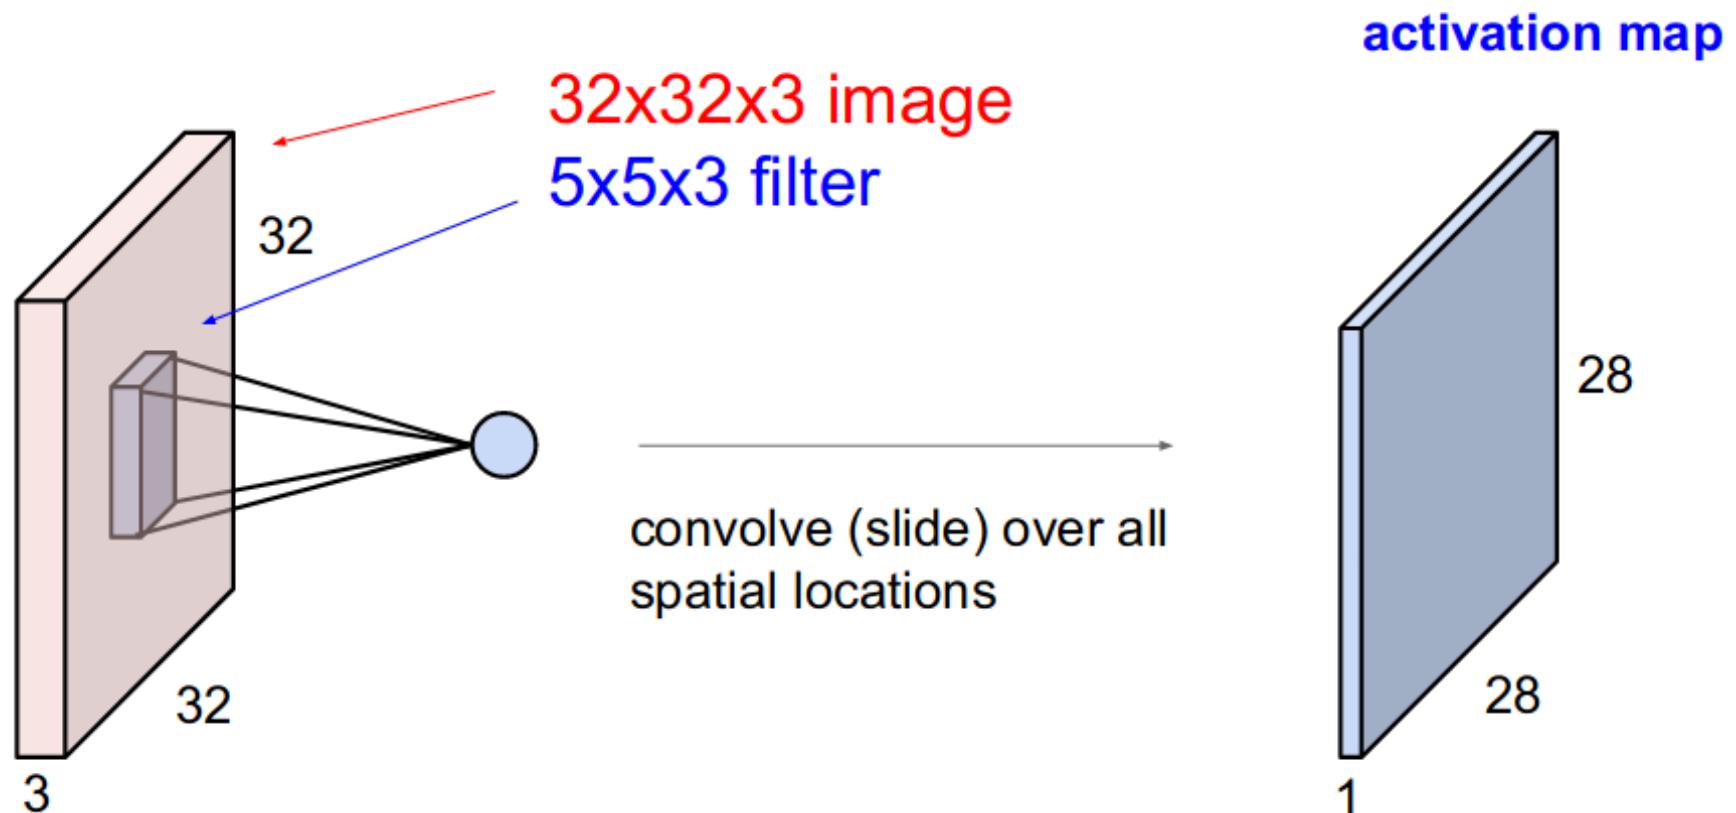


Fully-Connected Neural Network



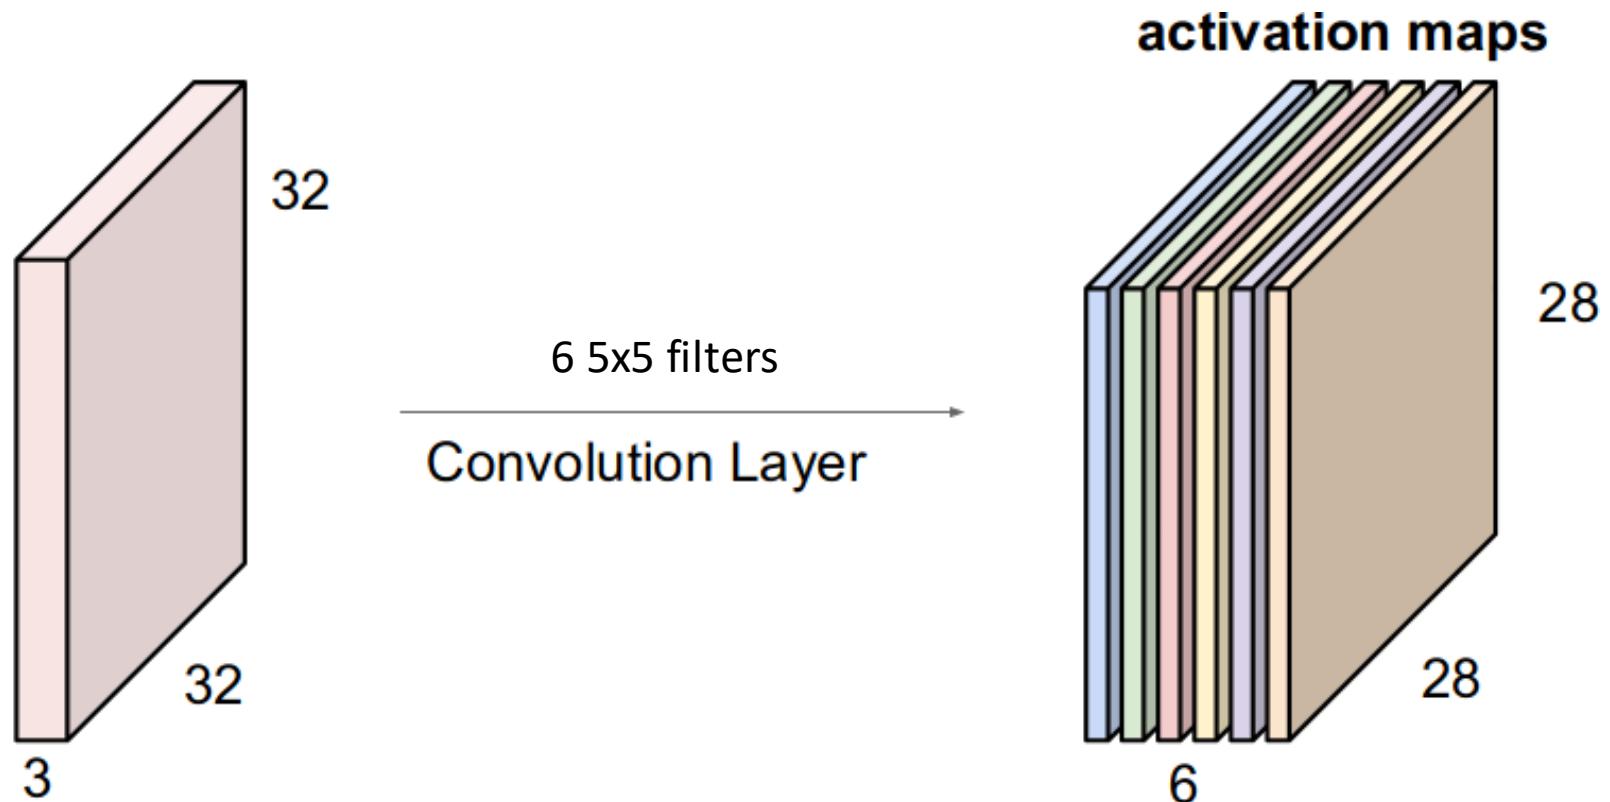
Convolutional Neural Network

Convolution Layer

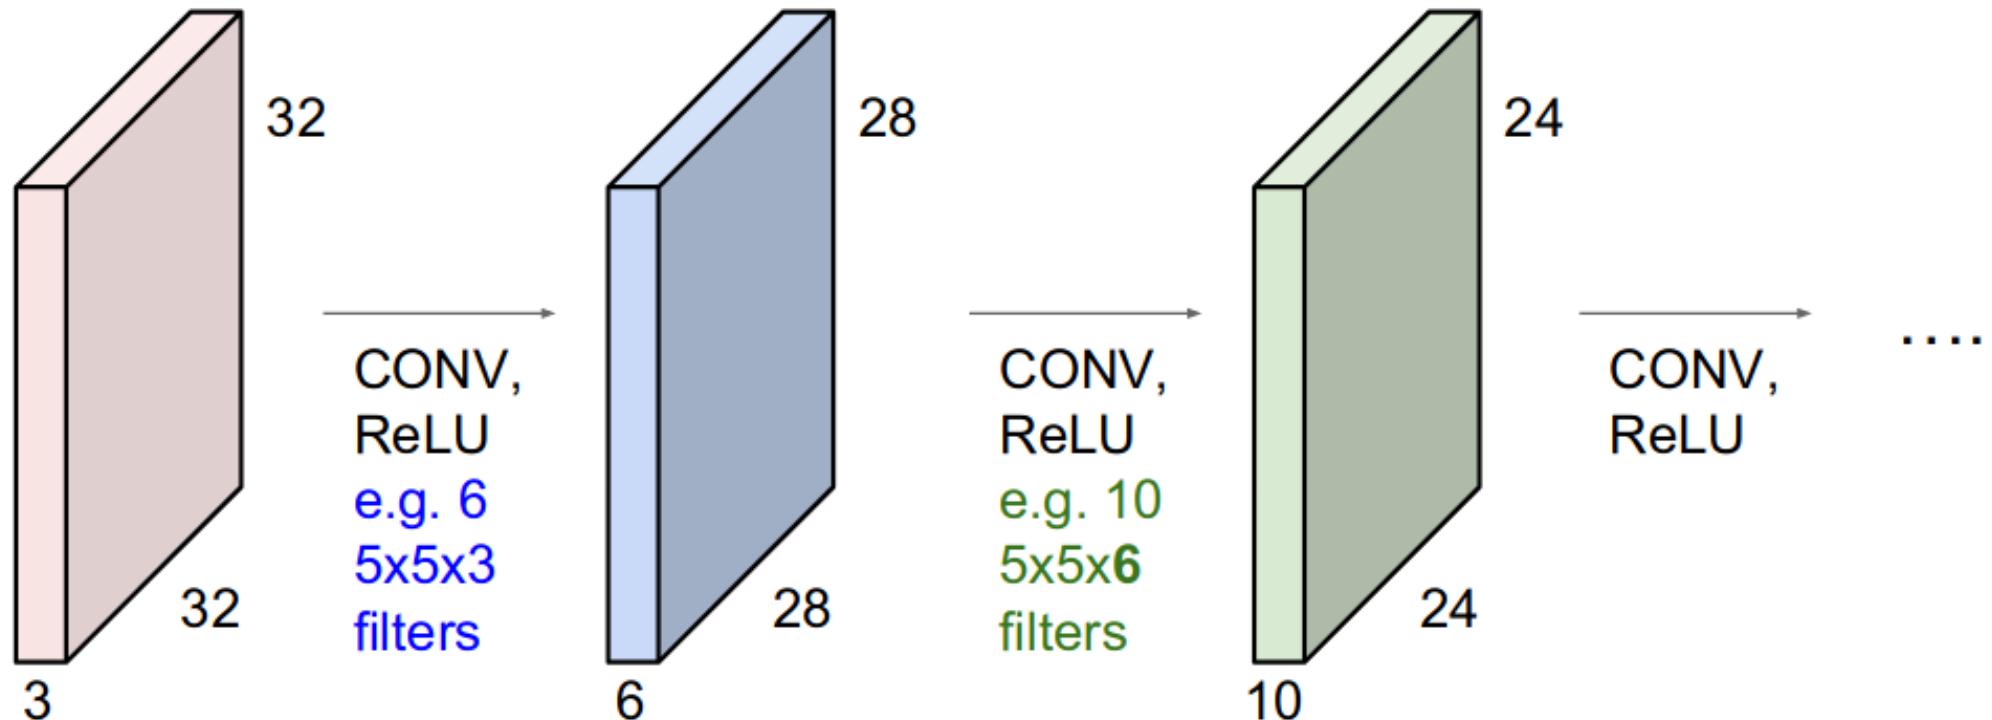


Convolve the filter with the image i.e. “slide over the image spatially, computing **dot products**”

Convolution Layer

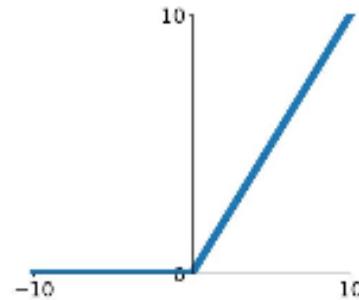


Convolution Layers



ReLU

$$\max(0, x)$$

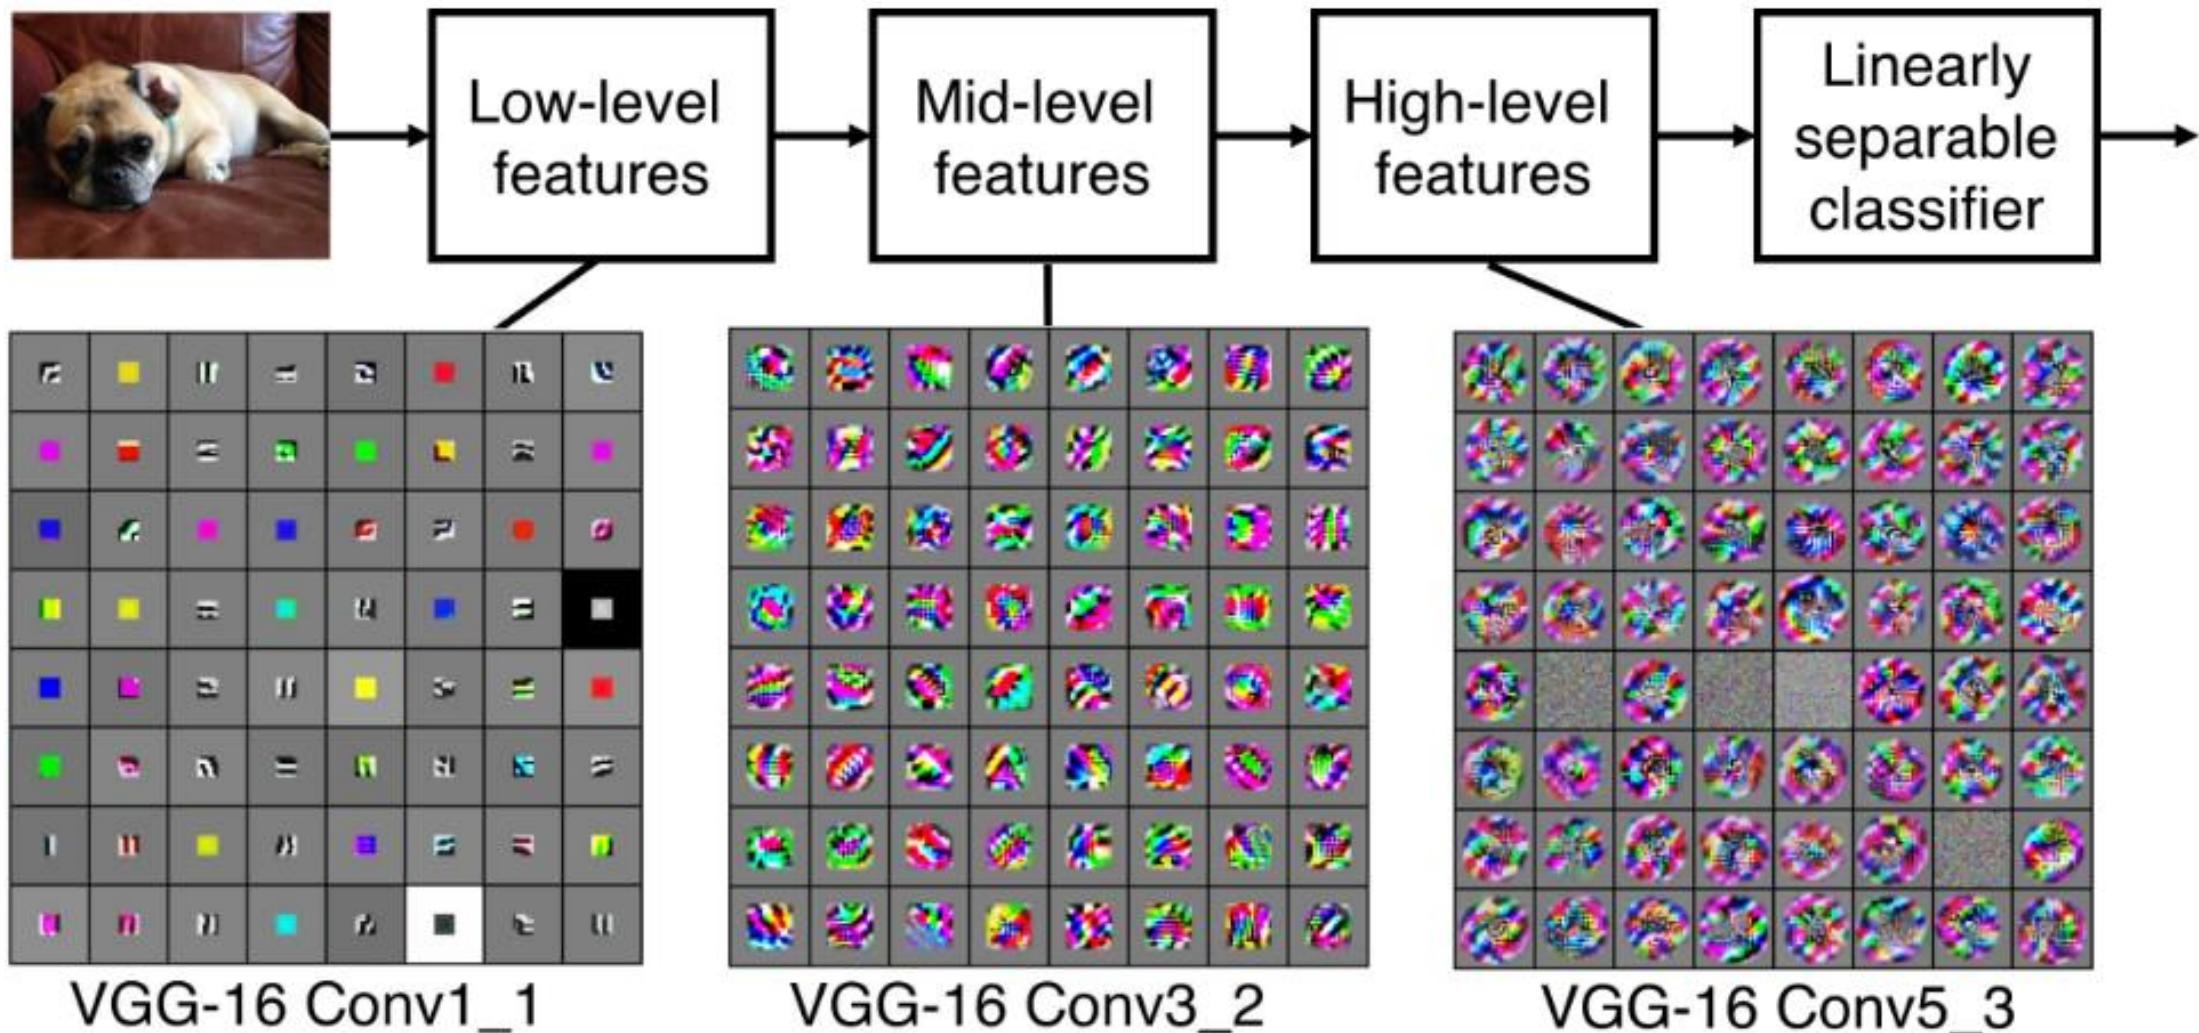


0	5	-1	2
-5	3	-4	6
6	4	0	-1
-1	8	-7	9

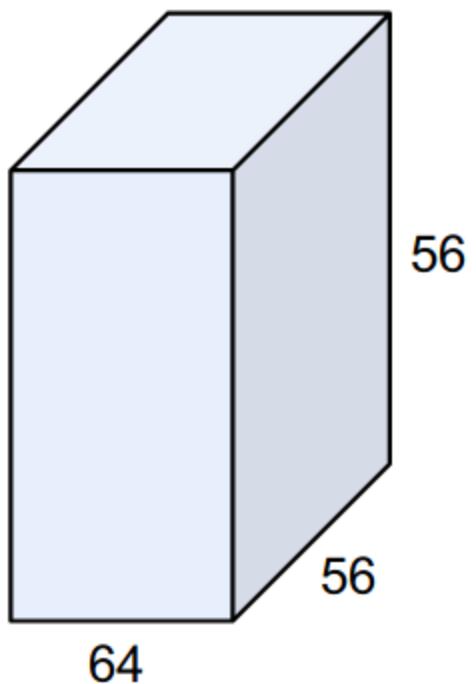
ReLU



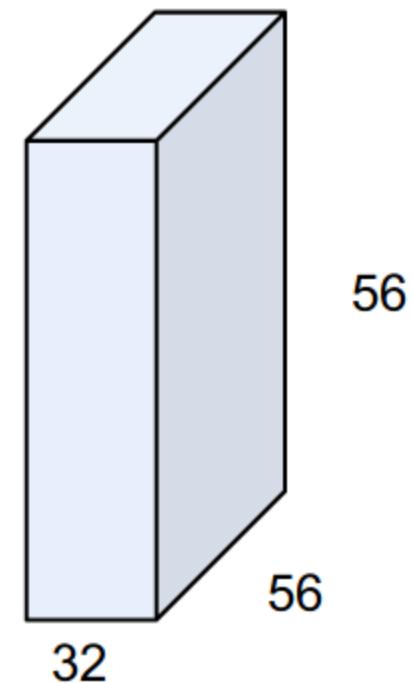
0	5	0	2
0	3	0	6
6	4	0	0
0	8	0	9



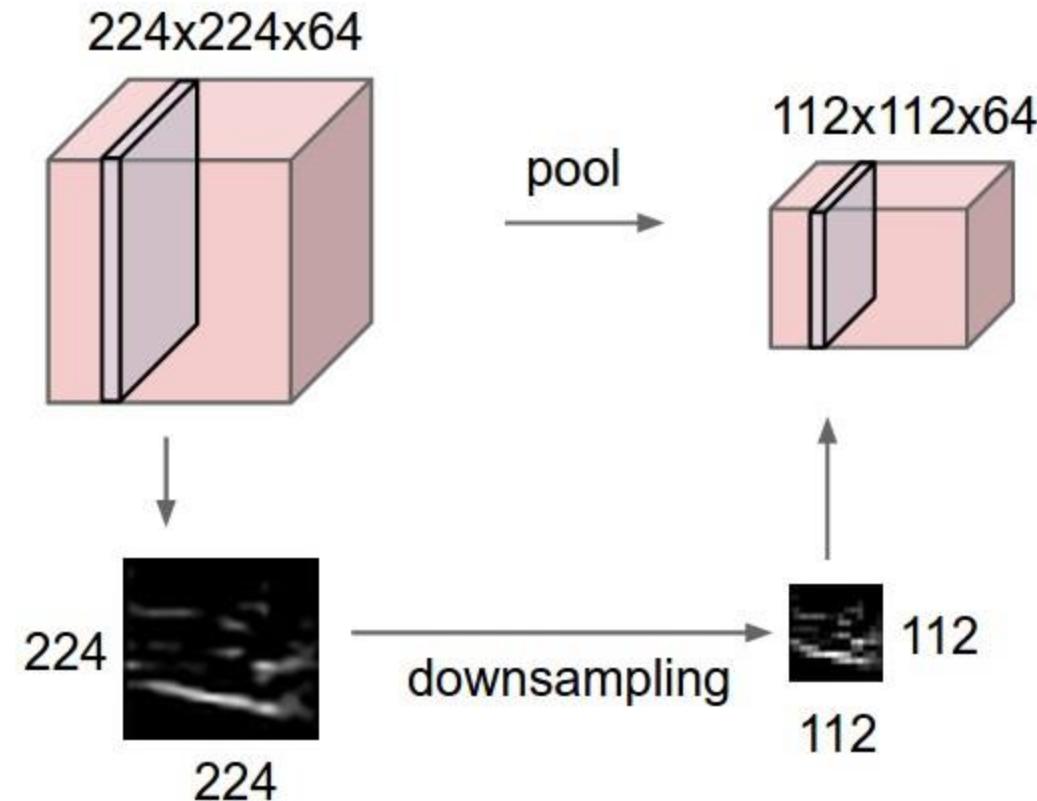
1x1 Convolution layer



1x1 CONV
with 32 filters
→
(each filter has size
1x1x64, and performs a
64-dimensional dot
product)

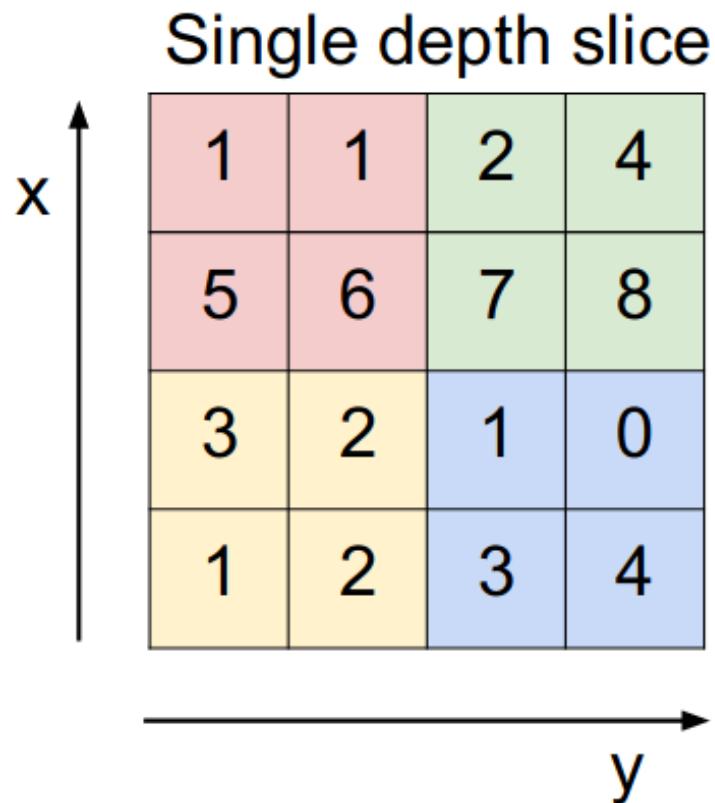


Pooling layer



- makes the representations smaller and more manageable
- operates over each activation map independently

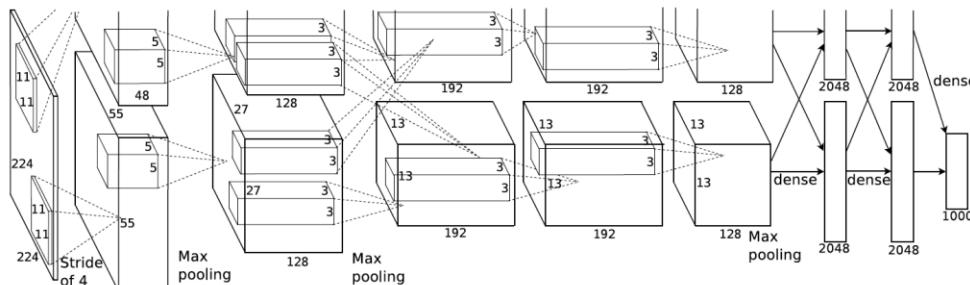
Max Pooling



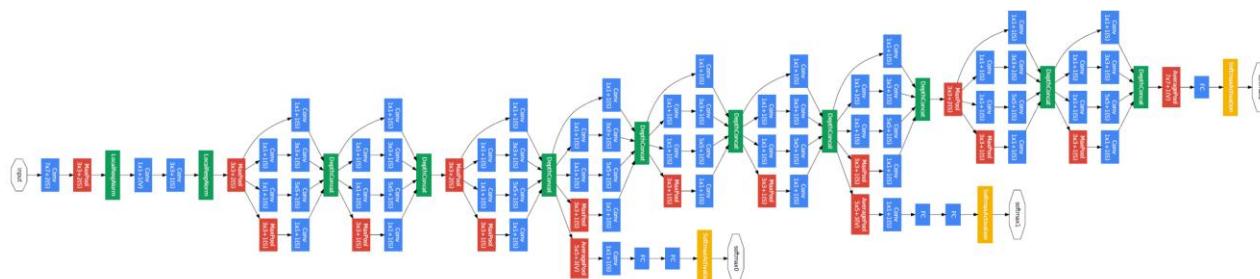
max pool with 2x2 filters
and stride 2

6	8
3	4

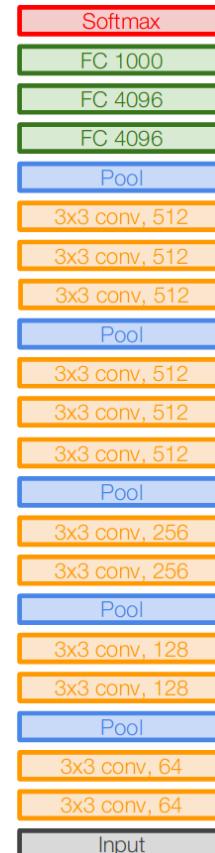
CNN Architectures



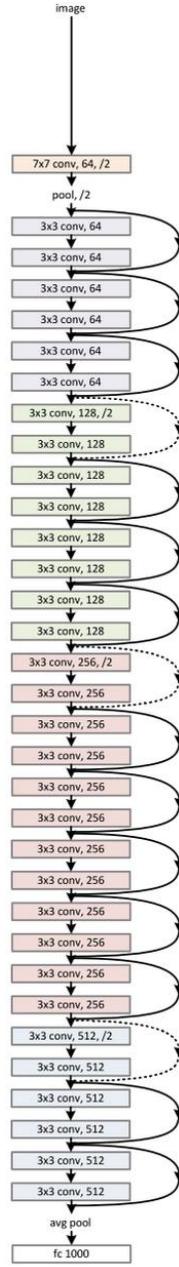
AlexNet



GoogleNet

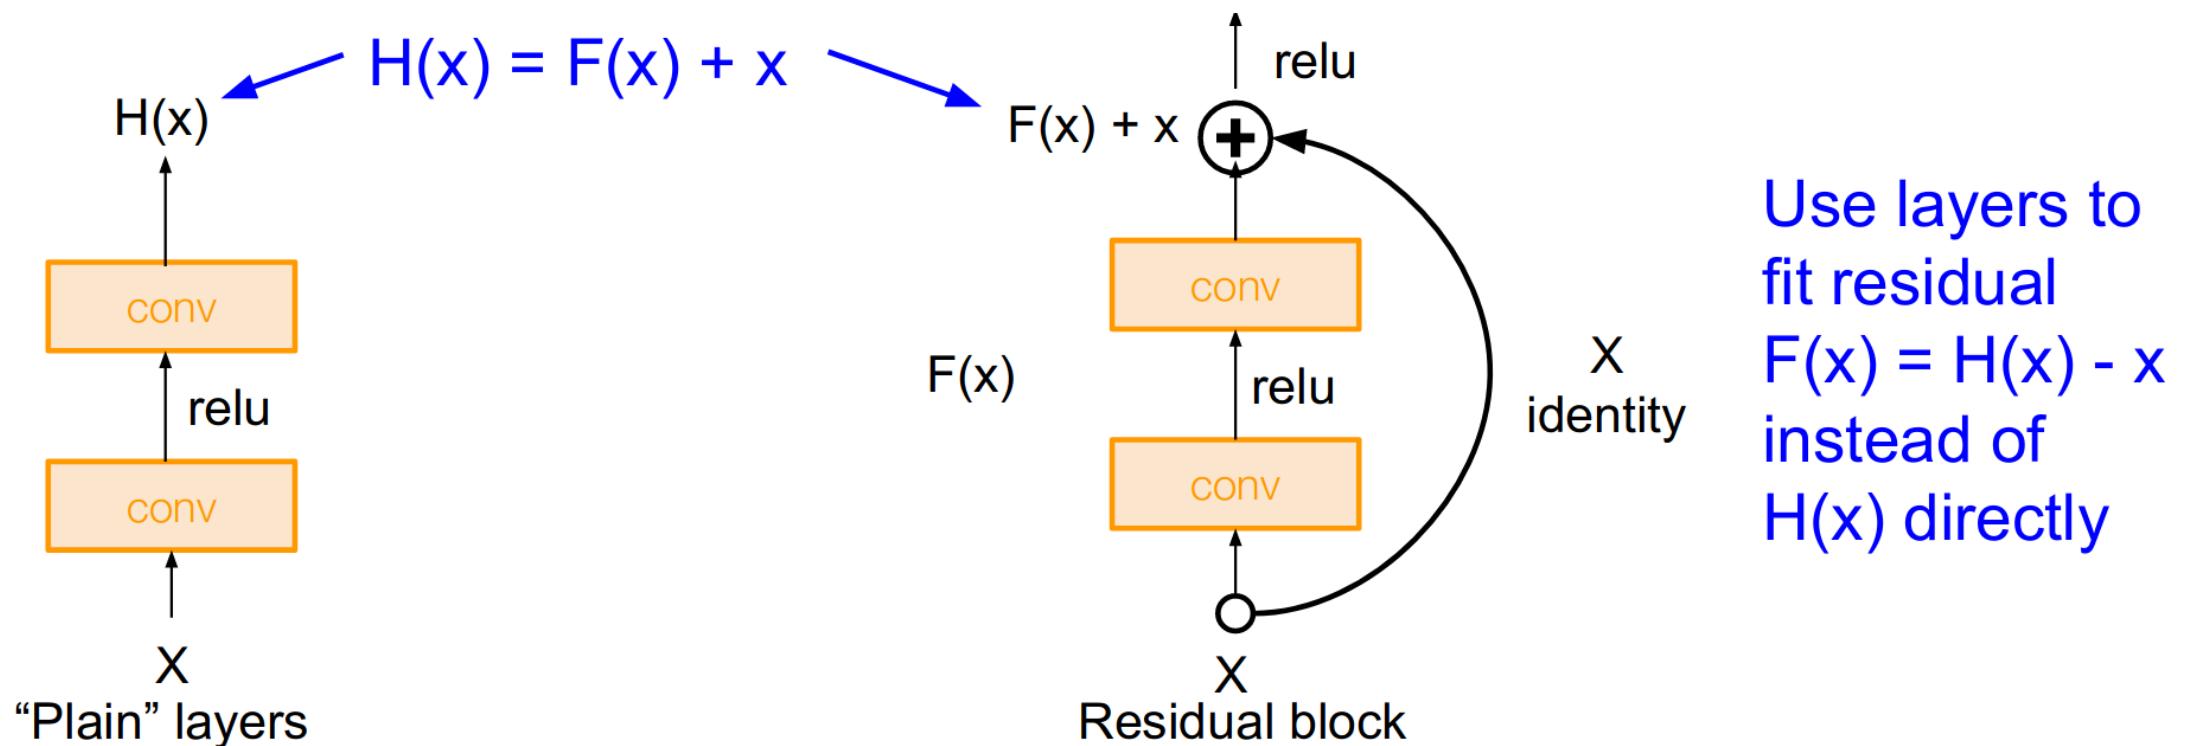


VGGNet



ResNet

ResNet



Loss Functions

Suppose: 3 training examples, 3 classes.

With some W the scores $f(x, W) = Wx$ are:



cat	3.2	1.3	2.2
car	5.1	4.9	2.5
frog	-1.7	2.0	-3.1

A **loss function** tells how good our current classifier is

Given a dataset of examples

$$\{(x_i, y_i)\}_{i=1}^N$$

Where x_i is image and
 y_i is (integer) label

Loss over the dataset is a sum of loss over examples:

$$L = \frac{1}{N} \sum_i L_i(f(x_i, W), y_i)$$

Softmax Classifier (Multinomial Logistic Regression)



cat	3.2
car	5.1
frog	-1.7

scores = unnormalized log probabilities of the classes.

$$P(Y = k | X = x_i) = \frac{e^{s_k}}{\sum_j e^{s_j}} \quad \text{where} \quad s = f(x_i; W)$$

Want to maximize the log likelihood, or (for a loss function) to minimize the negative log likelihood of the correct class:

$$L_i = -\log P(Y = y_i | X = x_i)$$

in summary: $L_i = -\log\left(\frac{e^{s_{y_i}}}{\sum_j e^{s_j}}\right)$

Softmax Classifier (Multinomial Logistic Regression)



Want to interpret raw classifier scores as **probabilities**

$$s = f(x_i; W)$$

$$P(Y = k|X = x_i) = \frac{e^{s_k}}{\sum_j e^{s_j}}$$

Softmax Function

cat
car
frog

3.2

5.1

-1.7

Unnormalized
log-probabilities / logits

Probabilities
must be ≥ 0

24.5

164.0

0.18

unnormalized
probabilities

exp

normalize

0.13

0.87

0.00

probabilities

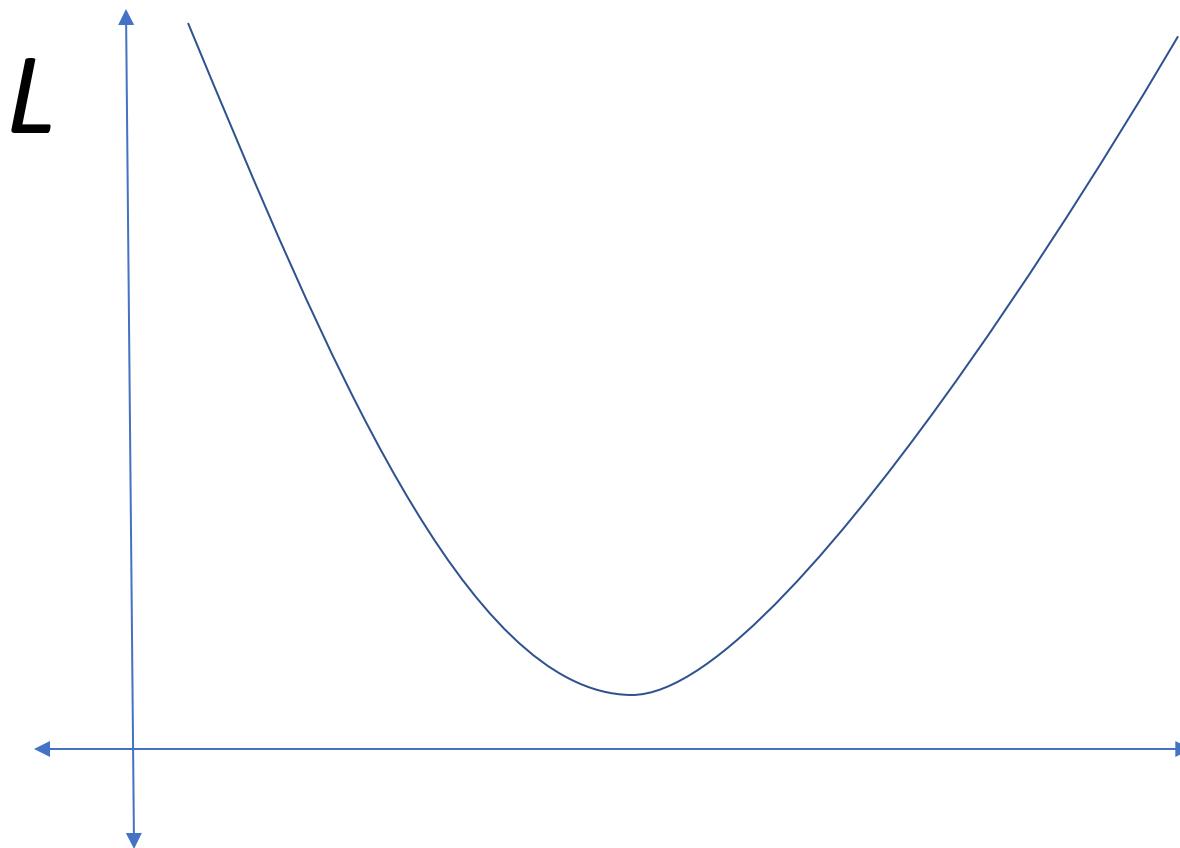
$$L_i = -\log P(Y = y_i|X = x_i)$$

$$\rightarrow L_i = -\log(0.13) \\ = 0.89$$

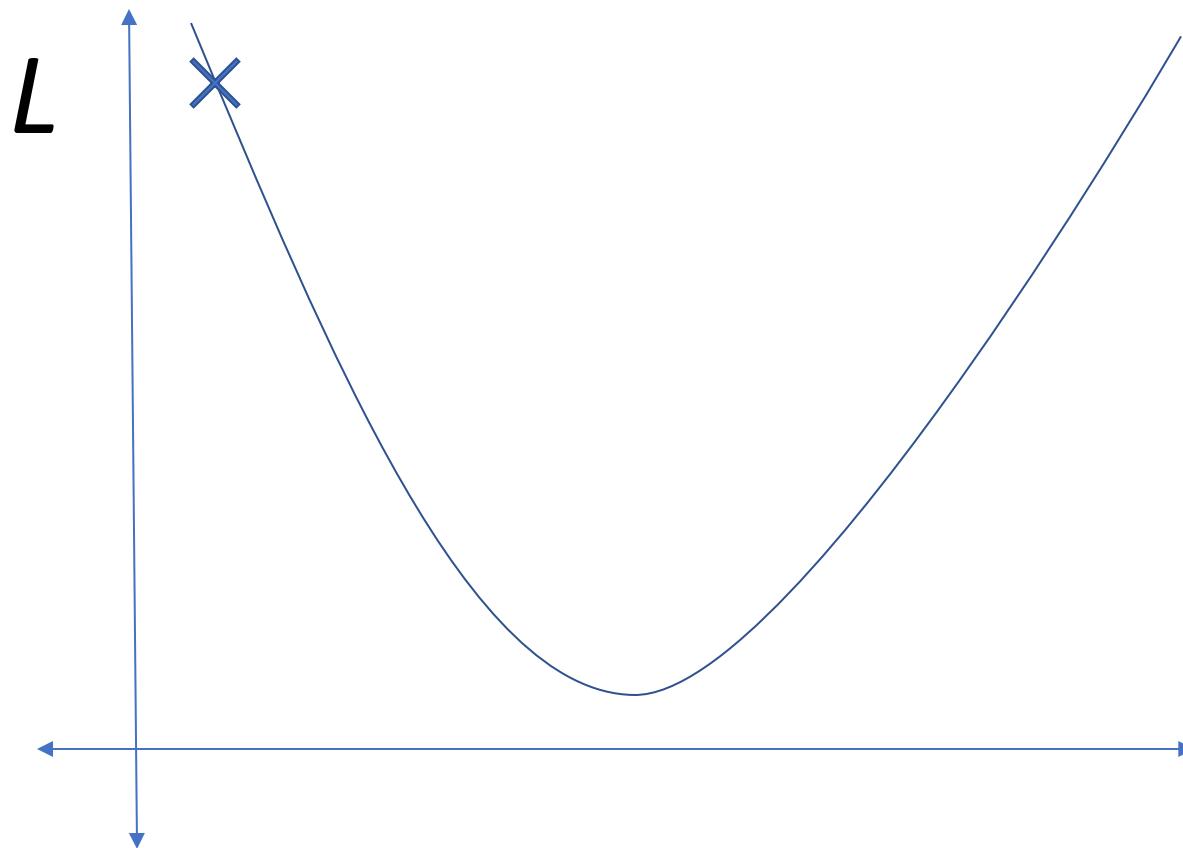
Optimization



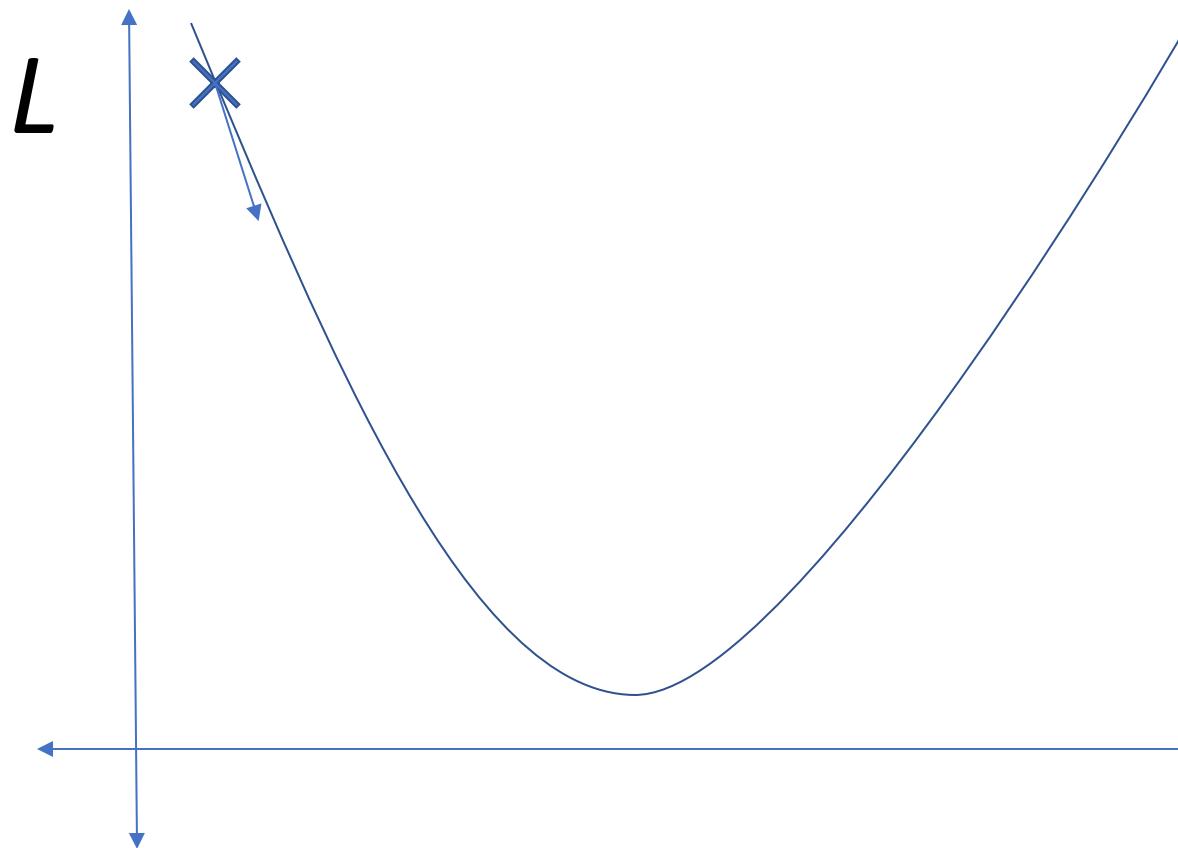
Optimization



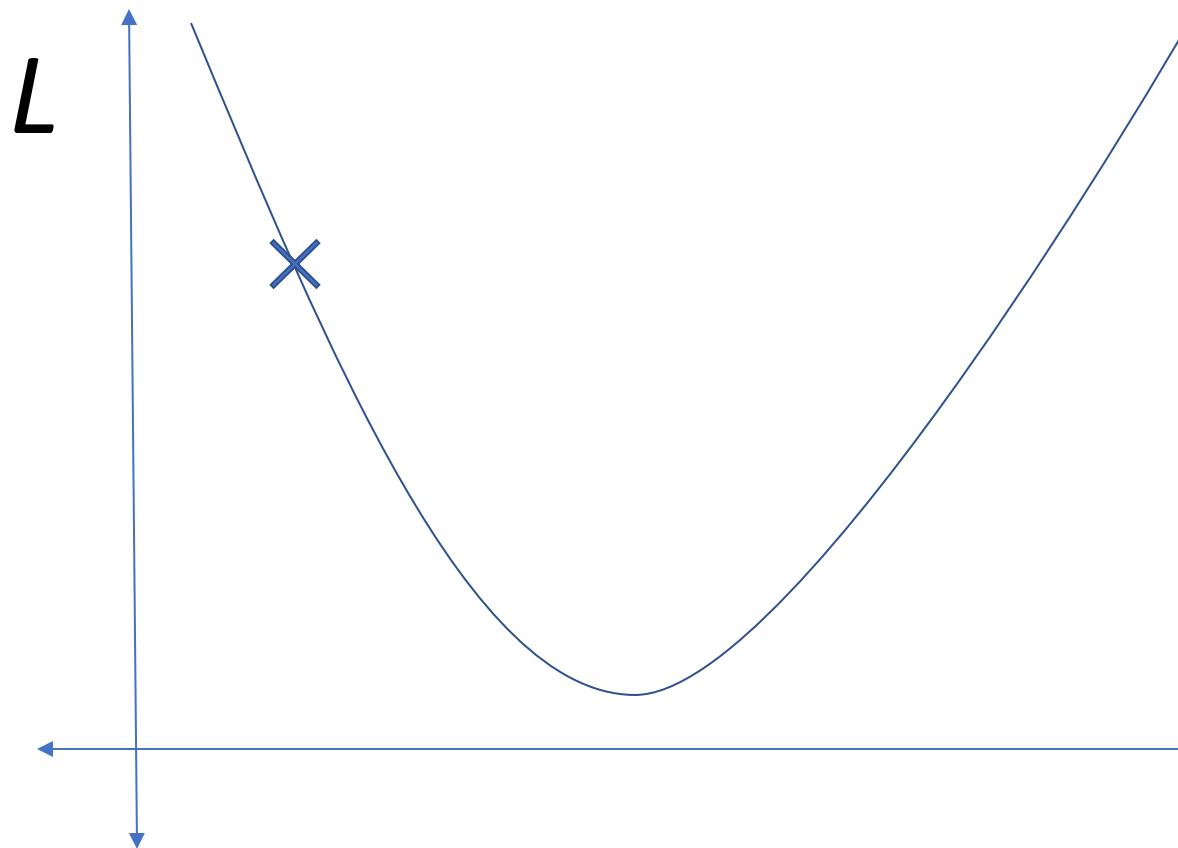
Optimization



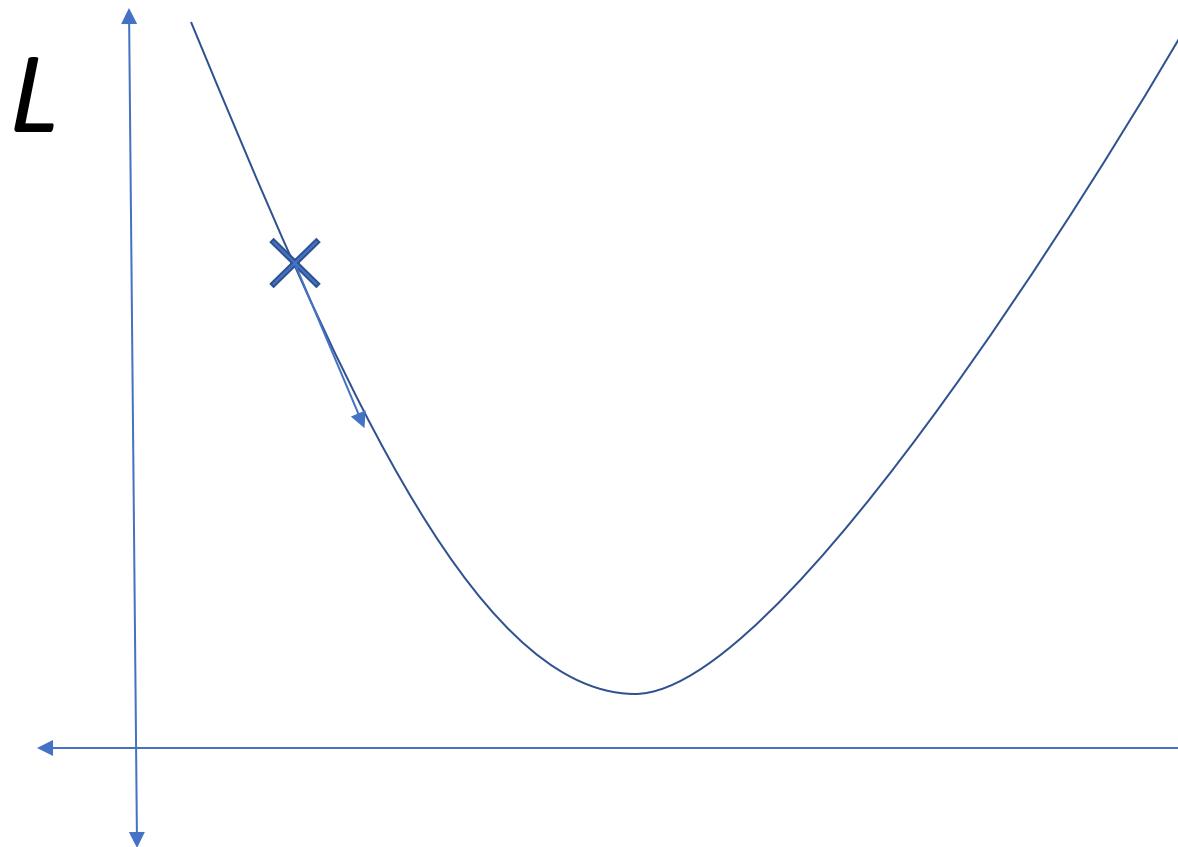
Optimization



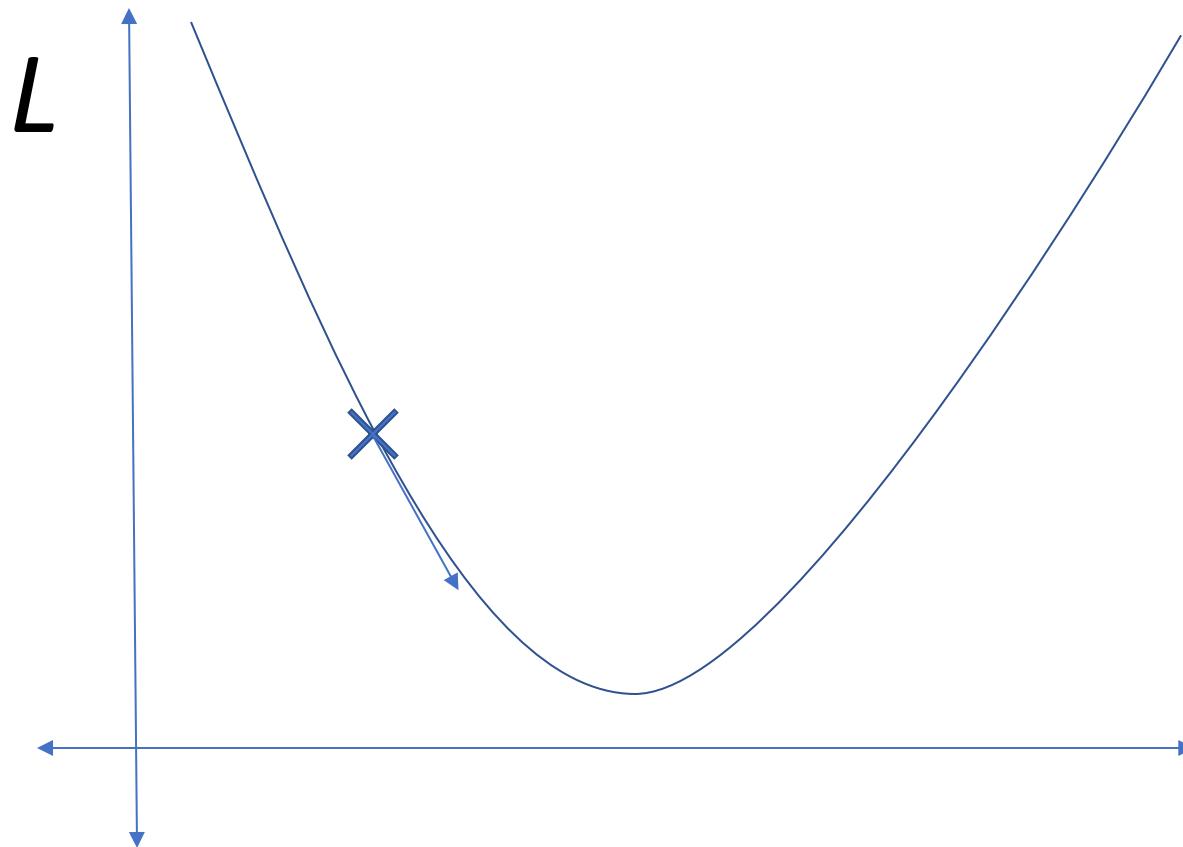
Optimization



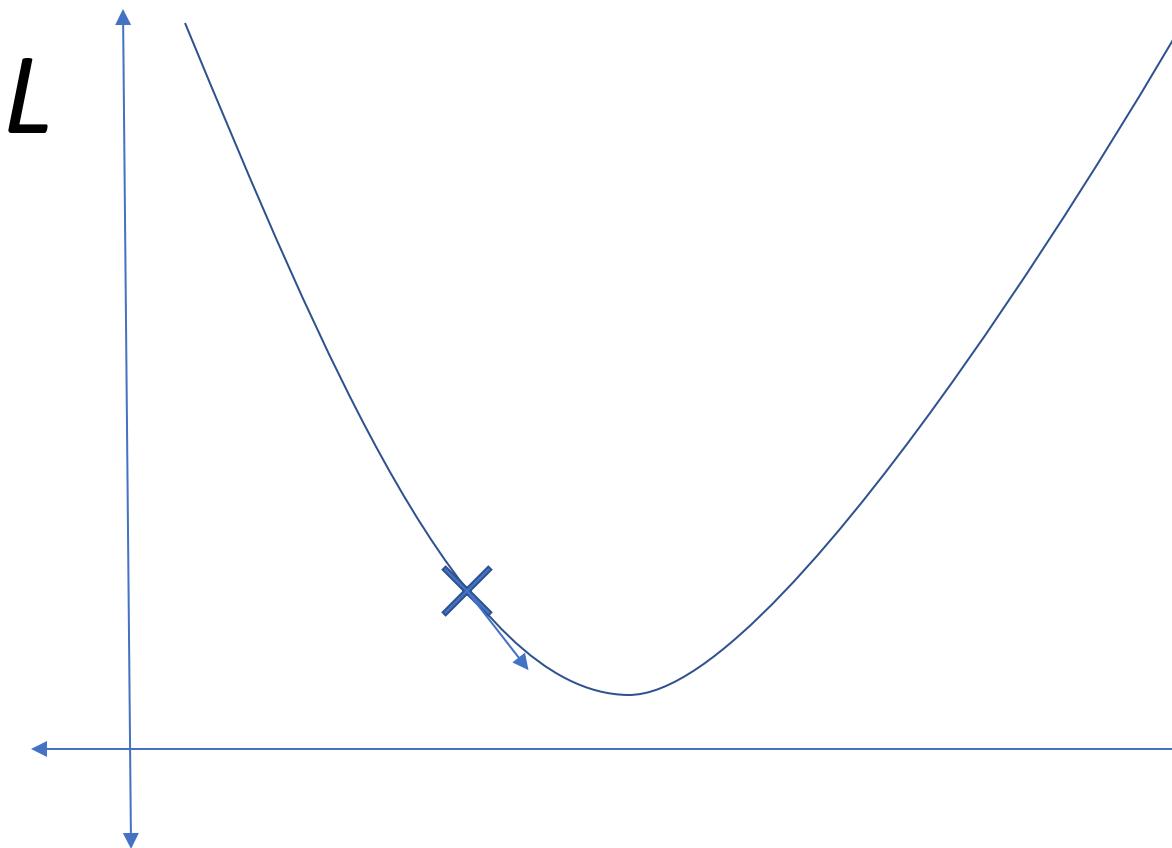
Optimization



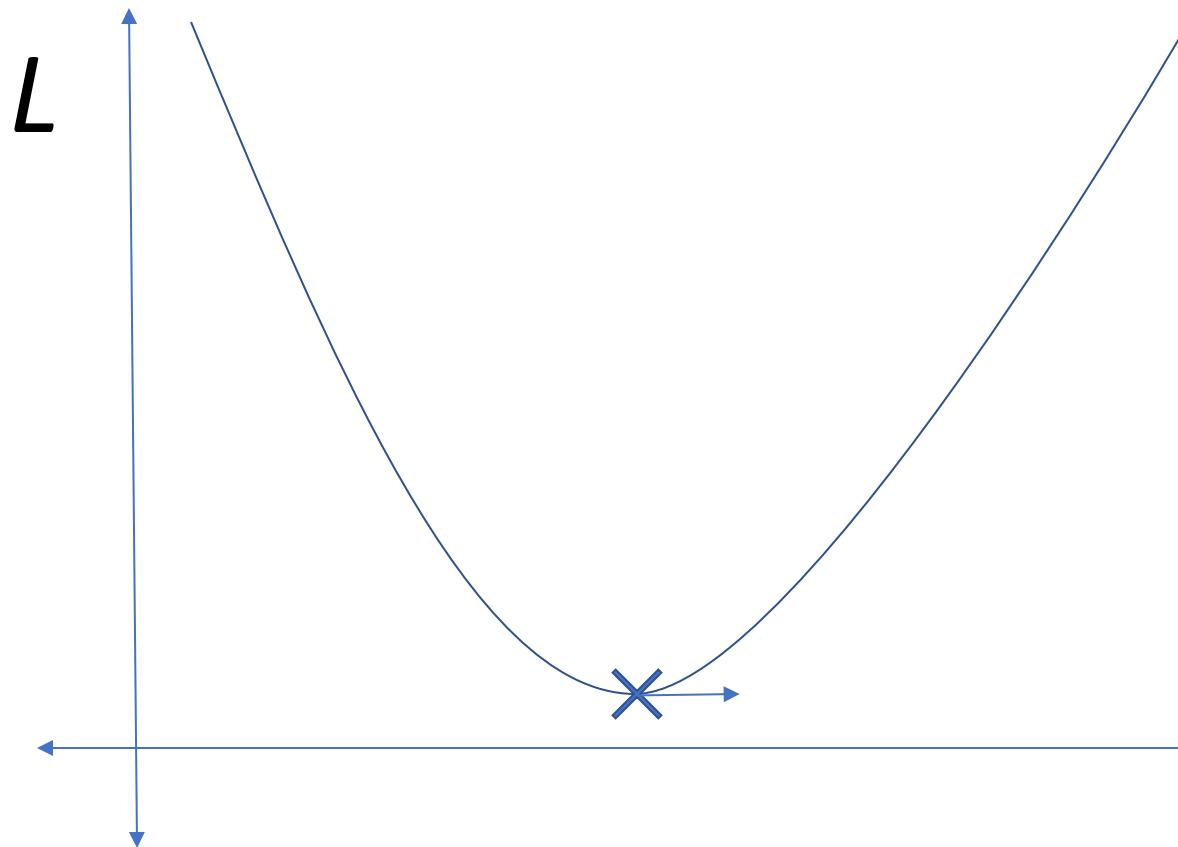
Optimization



Optimization



Optimization



Gradient Descent

$$L(W) = \frac{1}{N} \sum_{i=1}^N L_i(x_i, y_i, W) + \lambda R(W)$$

$R(W)$: Regularization term

$$\nabla_W L(W) = \frac{1}{N} \sum_{i=1}^N \nabla_W L_i(x_i, y_i, W) + \lambda \nabla_W R(W)$$

Gradient Descent

```
while True:  
    weights_grad = evaluate_gradient(loss_fun, data, weights)  
    weights += - step_size * weights_grad # perform parameter update
```

Stochastic Gradient Descent

$$L(W) = \frac{1}{N} \sum_{i=1}^N L_i(x_i, y_i, W) + \lambda R(W)$$

$$\nabla_W L(W) = \frac{1}{N} \sum_{i=1}^N \nabla_W L_i(x_i, y_i, W) + \lambda \nabla_W R(W)$$

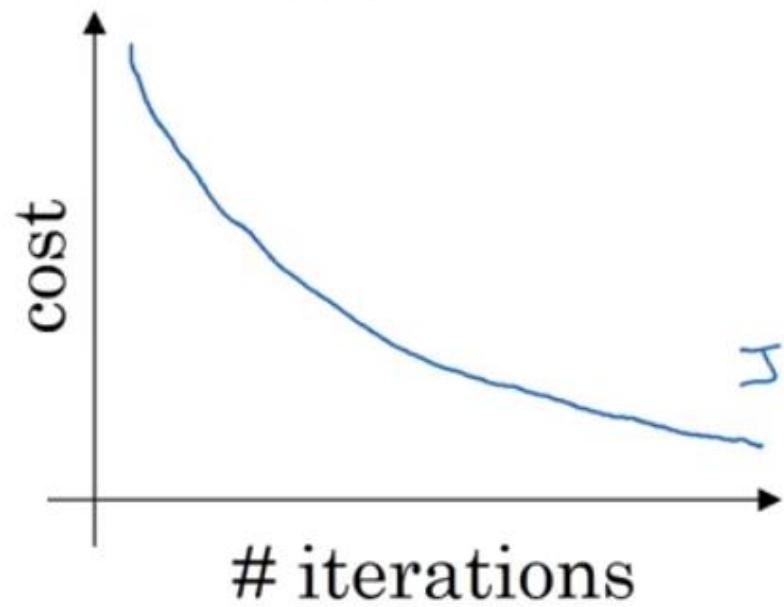
Full sum expensive
when N is large!

Approximate sum
using a **minibatch** of
examples
32 / 64 / 128 common

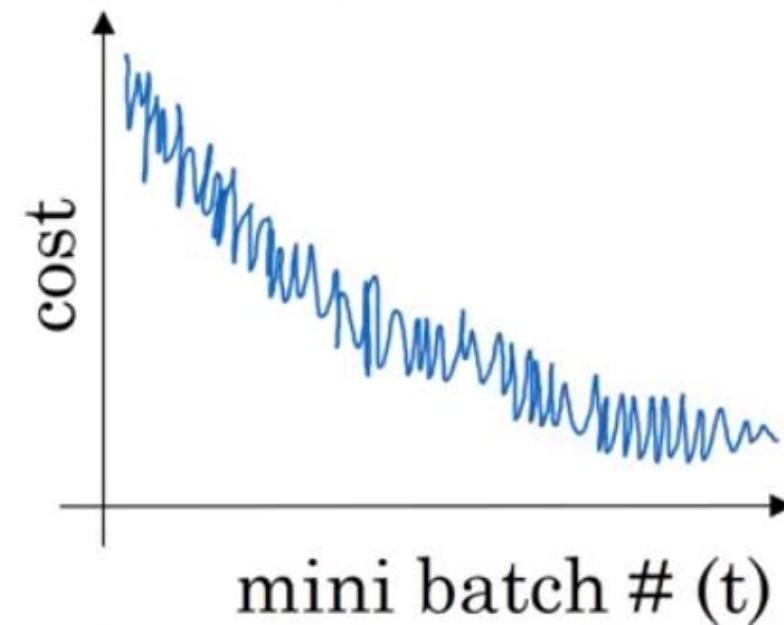
Stochastic Gradient Descent

```
while True:  
    data_batch = sample_training_data(data, 256) # sample 256 examples  
    weights_grad = evaluate_gradient(loss_fun, data_batch, weights)  
    weights += - step_size * weights_grad # perform parameter update
```

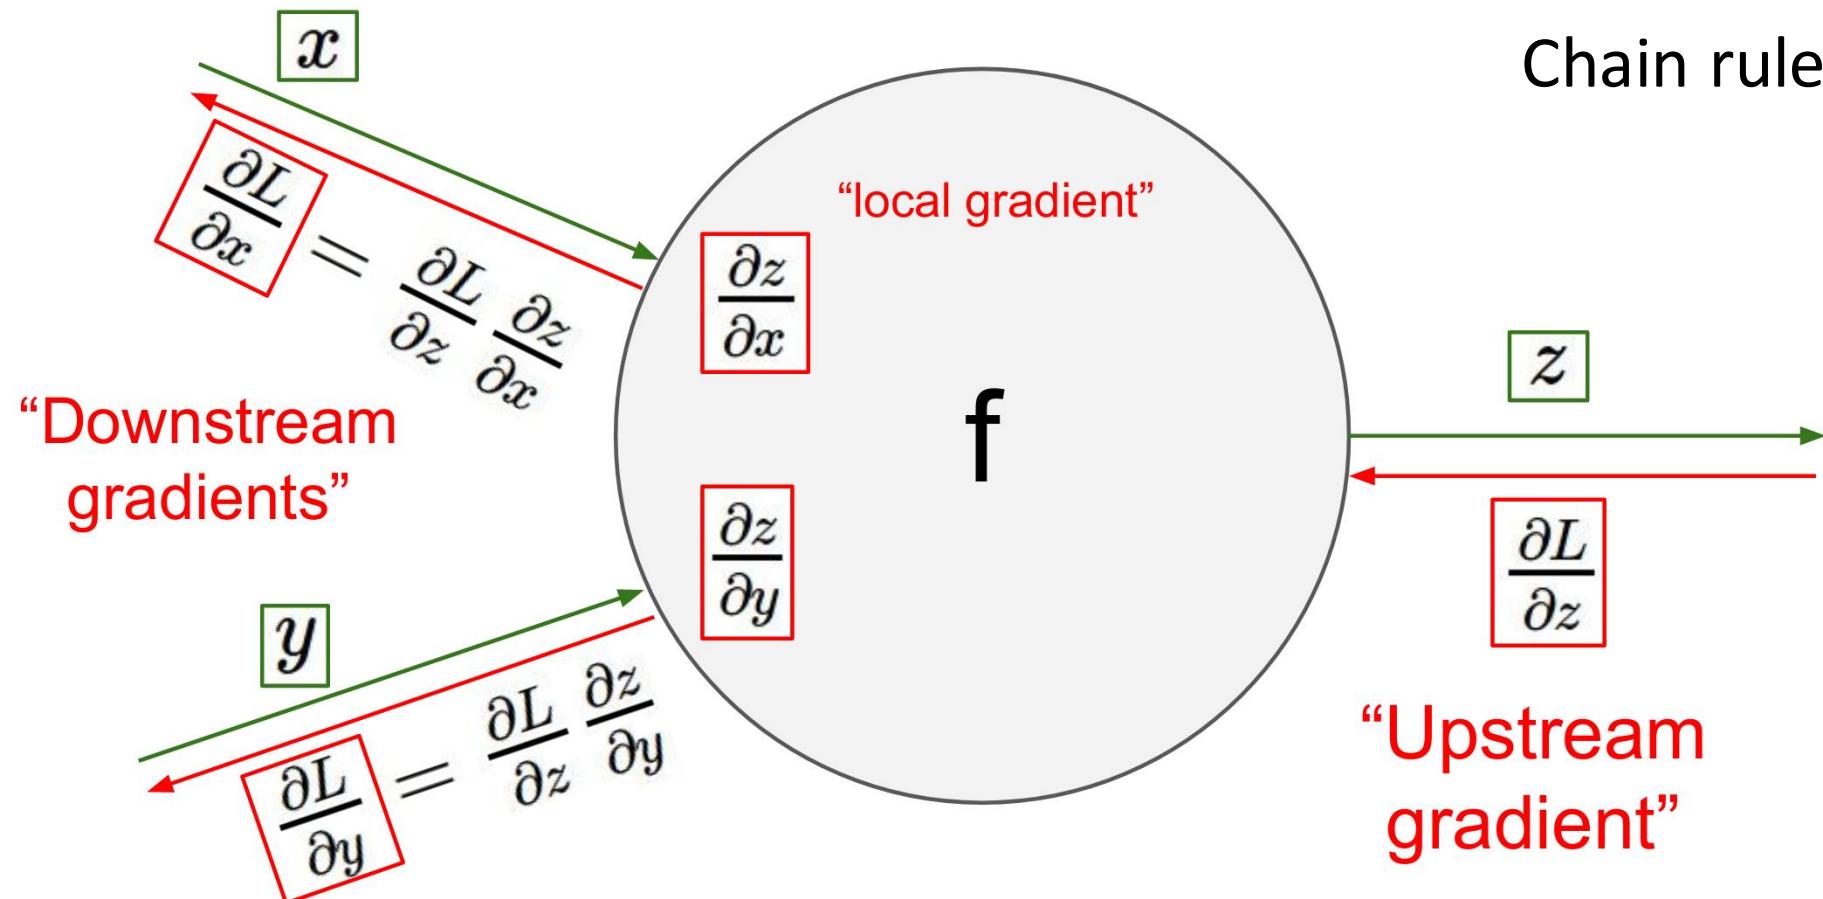
Batch gradient descent



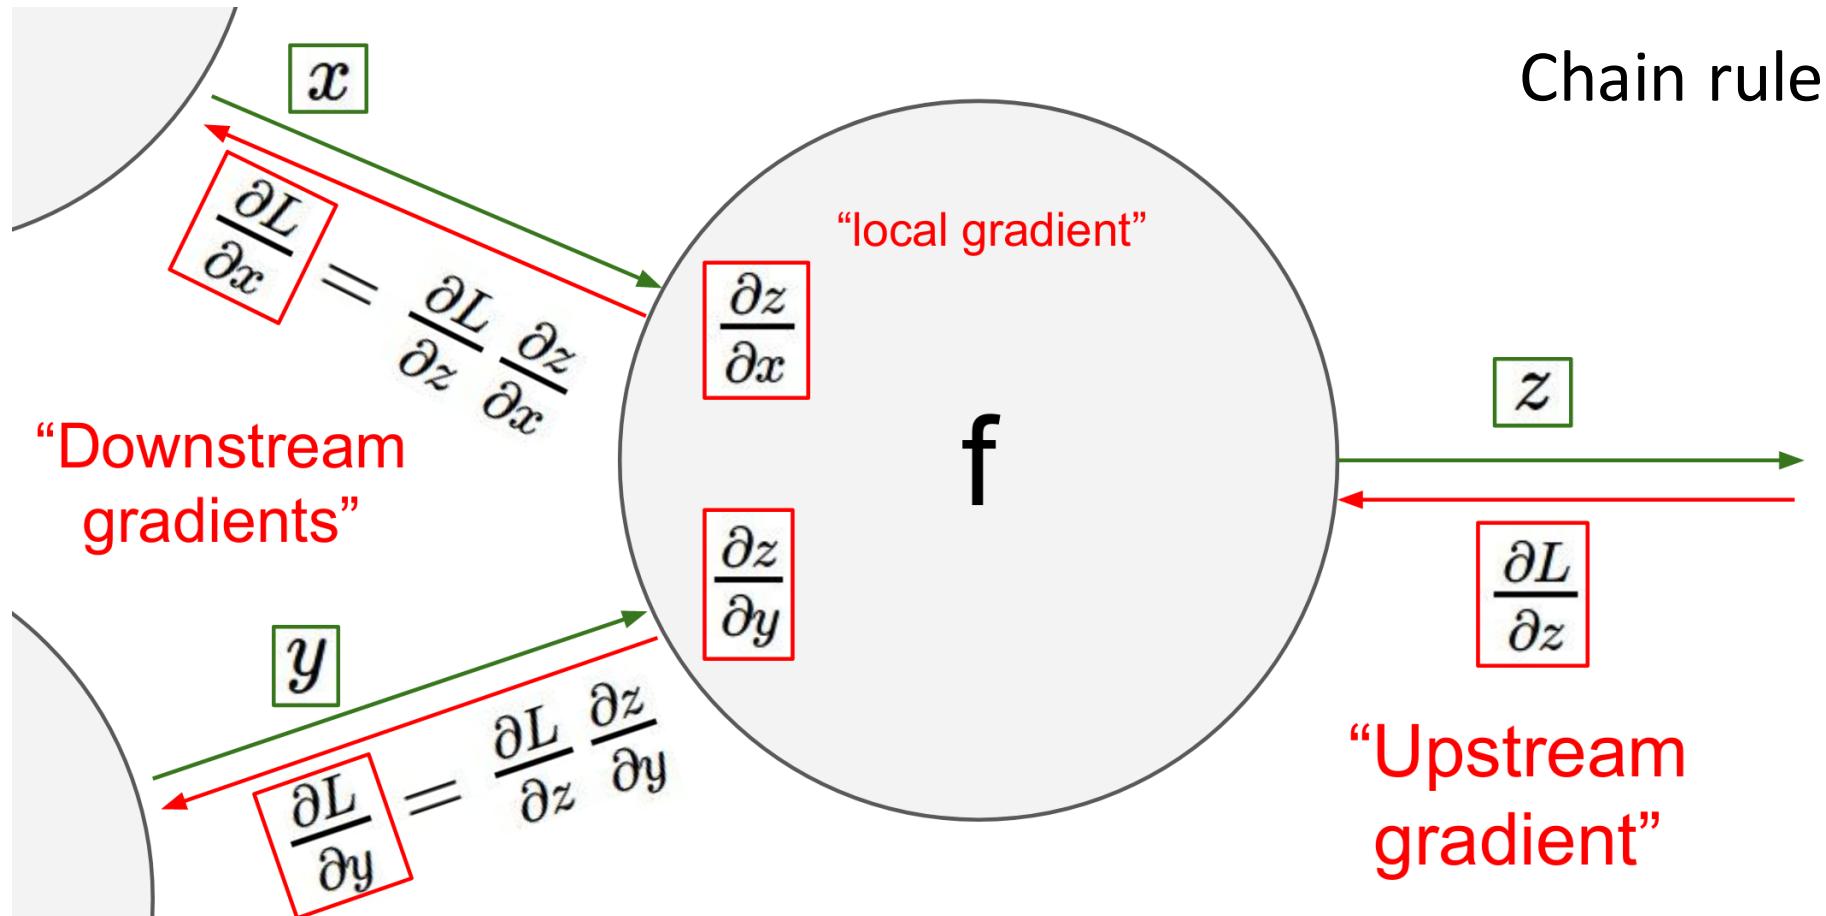
Mini-batch gradient descent



Backpropagation



Backpropagation



CNNs are Everywhere

